Weekly 3

Hazel Lutz

March 2, 2026

Weekly Progress (5 points)

1. Since I got an OutOfMemory Error on any device I tried running my datasets on, it appears that I must actually resort to doing the feature selection early. Oh well! It's done. I had to make the number of decisions about which columns to keep and what years to include, but since it's done, I guess I can let go of it. Just kidding! I'm going to go back and second guess all my decisions because that's important in social sciences.
2. Successfully took a number of other datasets done by the same organizations and merged them together (usually several waves/years of the same data collection) through for loops of filling a dataset. Some issues with this remain – column names changing over time for some reason I can only conclude is to make my life more difficult.
3. Tried on a few conflict event datasets as well with mixed success. Spent a while looking through those codebooks as well, which gets really rough when they're around a hundred pages long and very repetitive and dense. It's fiiiiine. But it does encourage me to reread them in the coming weeks to make sure I'm not missing something useful or misunderstanding variables.
4. Scheduled a meeting with a PhD student for this week to discuss a number of things, including methodology – though she specializes in different fields, it should still be useful for dealing with data.
5. Started some EDA in the R console before realizing I'm going to have to clean the data up a bit more before I start going hard on the EDA – something to work on in the next few days, I suppose. Mostly some of the aforementioned issues with columns and codebooks, but also some NAs I'll have to figure out how to deal with. It should all come out in the later analysis due to a lot of averaging based on conflict events.

Challenges (3 points)

1. The aforementioned problem with the large datasets was resolved! Lovely. Not that it's perfect – I did encounter some LFS issues while commiting this because even though I cleared up my like 1.5GB CSVs, a 150MB one remained. Darn. It did require feature selection I will continue to have to second-guess myself on and reread codebooks for.
2. This was also mentioned above, but some of the datasets remain inconsistent between and within themselves, and one of the conflict datasets I was counting on ended up not being what I needed upon closer examination – there are some conflicts I know it should have included but didn't, probably due to being too strict (or more strict than I'd say). I worry I'll have to do some manual work, but if I do, so be it.
3. The EDA revealed some more cleaning that needed to be done, so I suppose it's good that my early EDA process tends to be very improvisational and educational rather than interested in making a pretty product. That's for next week.

Code Commits (2 points)

1. Made the large CSVs manageable and produced a couple more large CSVs in the process out of other datasets: https://github.com/GW-datasci/26-spring-LUTZ-H/commit/fe1bb2b481d22be687a91ce7d6fbe47e335caaf1
2. Continued some more feature selection, narrowed down said large aggregated data, and started EDA: https://github.com/GW-datasci/26-spring-LUTZ-H/commits/main/