

Weekly 1

Hazel Lutz

February 9, 2026

Weekly Progress (5 points)

1. Spent a while acquainting myself further with git, especially git LFS (necessary given the size of some of the datasets I'm using). As the challenges demonstrate, this was far from a perfect experience. I intend to start anew with it next week rather than scrambling mid-commit.
2. Done some low level exploration of the data – public opinion data remains the hardest to deal with, but it also remains what I'm most passionate about as well as theoretically very relevant to statebuilding and peacebuilding.
3. Furthered quantitative research – was recommended a book, not on civil wars but on post-apartheid South Africa and the Truth and Reconciliation committee based on a very broad survey (very impressive!), and has thoroughly fleshed out the ways I consider the role of institutions in post-conflict statebuilding.
4. Definitionally, I believe I will be rather unopinionated on what constitutes a "civil war", though I will likely do some testing on different definitions of what exactly makes something a "civil war" to see whether conclusions will vary – this could provide some sort of secondary result for my research. My default usage of what constitutes a civil war will be based on the Correlates of War project, most likely, as this remains a fairly authoritative resource. This coming week's definitional hurdle will be on defining "success of recovery", which I will once again look to quantitative studies to give a decent consideration of.
5. Started my R work! Managed to get a bit of stuff done with it, despite the other, surrounding challenges of this week which has in some ways been less productive than desired. This is what I get for dealing with difficult datasets with no APIs, which I do miss dearly at this point. At the very least, I'm at the point where I can jump off tomorrow or the day after with slimming down the datasets I have into a more manageable format.

Challenges (3 points)

1. Though I'm used to simple tasks with git, having to deal with git LFS was a major hurdle. Lots of troubleshooting, lots of stackexchange, and I still haven't fully gotten where I need to be with it. For the sake of making any progress at all, I just chose not to commit any of the data. For legal reasons, it may be best if I just store the data outside of the repository.
2. My early attempts at importing and EDA have gone horribly awry, as the datasets are all *very* different, not only between organizations, but also within results from the same organization. I will always praise good data management. Public opinion data is not where one finds such things. Trying to import AmericasBarometer data crashed my computer. I will be finding new avenues to parse this, more than likely in a more efficient language (likely R).
3. Having to deal with codebooks is a unique kind of pain. The start of this project is looking to be rather tedious! But that's okay. Having to deal with a 1.8GB or so .dta file means that I'll have to

slim it down quite a bit before I do really anything with it, as even in R it did crash pretty quickly (but not as quickly as Python, so I did get a colnames out of it!). Since the actual feature selection for the quantitative analysis is a bit down the line, I'll have to do a very broadly considered feature selection to narrow some of these files down, but that's okay.

Code Commits (2 points)

1. Early banging my head against a wall with data wrangling and cleaning and aggregating: <https://github.com/GW-datasci/26-spring-LUTZ-H/commit/10776674bd5efb4f51f3f5ef2b7cbe723d0dc88a>
2. Started R imports and cleaning:
<https://github.com/GW-datasci/26-spring-LUTZ-H/commit/ec8d2c6f1619f4f895342983dc51553b3e30edb6>